

**Predictive Recommendation Engine
Using MovieLens Dataset**

Final Project
Milestone #3

By:

Timur Zambalayev
Joshua Coffie

Harvard University

CS-109A, Fall 2016
Introduction to Data Science

Before beginning analysis and building a model on a dataset, we must first get a sense of what we're working with. We've found that datasets can often be skewed, based upon observation method and a myriad of other factors, so gathering initial statistical data is the first step in building any form of predictive model. Our approach is to attack our initial analysis in three parts - the first is an overall summary of the data, the second is analysis from the viewpoint of the reviewers, and the third is from the movie perspective. After this initial analysis, we fully expect the data to lead us into additional areas of exploration.

After importing the data from three separate datasets and joining those data frames by movieID, we can see the kind of data that we're working with. The MovieLens dataset has just over 100,000 ratings on 9,100 movies from 671 unique reviewers. So, on the whole, we should be looking at around 10 reviews per movie with an average of 150 reviews per user (initial estimate). This is a scaled version of the entire dataset available from MovieLens that is designed for projects such as ours. It is important to note that we expect our project results, using this dataset, to hold even with additional observations. We expect to use the larger dataset to test our model by randomly selecting user data and evaluating the model's performance.




Figure 1: Determining the mean rating

Each observation in our dataset contains a variety of information: a unique movieID (an integer value) to identify the movie in question, the title of the movie (a string with the year of production on the end), associated genres for the movie (a string of values separated by a '|', in which there is a one-to-many relationship - one movie may have multiple genres), a unique userID (an integer that identifies a reviewer), a rating of the movie by the reviewer (a float which can range from a 0 to 5.0, within .5 increments), a timestamp for the review (ISO8601 based, milliseconds since epoch format), and tags (which consists of a string that a reviewer might use to describe the movie or their review).

Cleaning the data from the source was simpler than we expected. Since our dataset didn't have dive into lengthy reviews for each movie but captured user ratings, genre, and a short tag to describe the reviewer's take on the film, we were fairly fortunate to not need a lengthy cleansing process. Some films didn't receive a genre tag and were subsequently tagged as 'no genres listed.' In this case, the data was temporarily converted into an empty string. As we progress throughout the project, we will breakout the tagged genres into a predictive feature for each

film (which will include removing the '|' from each string) and use those responses to build our model. For the purposes of initial analysis, the string was replaced with a blank value. Otherwise, the dataset was particularly well-structured and documented, readily available in an easily consumable CSV format.

When exploring data, it's always best to get a sense of the range, mean, and frequency of values at the very beginning to determine if there is any readily apparent skewing of results or obvious trends. We did this both from the perspective of the movies (computing and visualizing average ratings, the frequency of rating, etc.) as well as from the reviewer's perspective (reviewer's average ratings, highest average ratings, frequency of ratings, etc.). We found that the mean rating for a movie by a reviewer was 3.54. We can use the closest value to the mean (3.5, or rather, 3.5/5.0) as our simplest (naive) baseline prediction. This mean will be used to sanity check our models in the future.




Figure 2: Mean ratings for Reviewers

Looking at the data from the reviewer's perspective, we find out some interesting information about the dataset; there appears to be large range for the number of ratings that a user might submit. For instance, the reviewer who contributed the least amount of feedback rated 20 movies while the most active reviewer submitted nearly 2400 ratings, whose individual contribution accounts for approximately 2.5% of the total reviews. We also see that some users are more prone to leave positive reviews while other users are quite the opposite; the highest reviewer's mean rating is 4.95/5.0 while the lowest reviewer's mean rating came in at 1.33/5.0.

The next step is to analyze how ratings apply to the films in question. Performing simple analysis on our dataset, we find that every movie has at least one review and the mean number of reviews per movie is about 11. We also see that the mean rating for a movie is 3.29, which is lower than the mean rating that a reviewer submitted. This is evidenced when we view a histogram of movie mean ratings and find spikes appearing at the .5 increments. This is because so many of our movies in the dataset only have 1 or 2 reviews.




Figure 3: Movie Mean Ratings

This leads us into another point of interest. Unfortunately, nearly 50% of the movies only have between 1 and 2 reviews. This is particularly concerning because these movies may be rated extremely high (or low) by one or two reviewers, which could affect how well our recommendation model works if we do not take this into account. Because of this realization, we fully expect to introduce prior beliefs into our model to adjust. For instance, when we recommend a similar movie to a user, we could return a short list of films instead of populating every movie that might qualify, especially if films might only have 1 to 2 reviews.

Using a log scale to adjust for a less skewed distribution, we do find another interesting aspect of the data when we compare the number of ratings per user and the user mean rating. We find that the more ratings we have from a user (i.e., greater than 10 ratings), the higher the likelihood that their average rating will gravitate toward our mean (between 3.0 and 3.5).




Figure 4: Number of Ratings per Reviewer Compared to Average Rating for a Movie

Using a similar comparison for the number of ratings per movie and the mean movie rating, we find that the more ratings we have per movie, the more likely the rating for the movie will be higher. This suggests that reviews are more prone to review movies they like than those they didn't like.




Figure 5: Comparing Number of Reviews for a Movie to Average Rating

When reviewing the dataset, we found some other areas that we wanted to explore. The first was the range of dates for reviews. The theory is that when a movie was watched and subsequently reviewed might have an impact on the rating. For this reason, we plotted the timestamps of the reviews in a histogram to see when they were submitted. It appears that the ratings were submitted cyclically and that the highest frequency of submissions occurred near the later part of 1999. Although this was not enough to establish our theory, we could take this visualization and compare it to the production years of the films to determine if it had an impact on the rating.




Figure 6: Distribution of Timestamps for Ratings

The last area we explored pertained to the genres of the films. We found that there were 19 distinct genres accounted for in the movies (with a 20th originally determined to be 'no genres listed'), the maximum number of attributed genres for a film came in at 10 (which seems pretty high), 18 films were not tagged with any genre, and the mean number of genres per film was 2 and 1/4. The mean seems to make sense as many films are a combination of a few genres, as this helps the movie appeal to more viewers.




Figure 7: Number of Genres per Movie

We determined that the most popular drama, both in number of movies and number of ratings, was drama, followed closely by comedy. Following this same train of thought, we discovered an interesting relationship among the third and fourth place genres; action is in third place for number of ratings, but it's 4th in the number of movies per genre, falling behind thrillers. In essence, action movies are more popular (have more ratings) but there are more thriller movies available for customers.




Figure 8: Mean Rating per Genre

Another interesting find involved documentaries, war, and film-noir genre films. Each of these genres is on the lower end of the spectrum when it comes to the number of ratings per genre and the number of movies per genre, but we find that they rank right near the top in average rating per genre. We suspect this is because these genres tend to attract special interest groups that are particularly supportive of their preferences, and may feel more inclined to highly review these genres. This is important to note - if one of these genres is highly liked by a user, there may be a good chance that they're eager to see another film in this space again. We imagine this is another space that prior beliefs may adjust our model to be a better fit.

NOTE: The remainder of our visualizations are available within our Python Jupyter Notebook on our repo. Please visit this site to view the extent of the analysis, outside of this summary submission. Thanks! Github repo:

https://github.com/starbuck10/CS109a_DataScience_UserRatings_Team_Project

Milestone #3 Folder:

https://github.com/starbuck10/CS109a_DataScience_UserRatings_Team_Project/tree/master/Milestone%20%233

Jupyter Notebook direct link:

[https://github.com/starbuck10/CS109a_DataScience_UserRatings_Team_Project/blob/master/Milestone %233/MovieLens/User_Ratings_Project.ipynb](https://github.com/starbuck10/CS109a_DataScience_UserRatings_Team_Project/blob/master/Milestone%20%233/MovieLens/User_Ratings_Project.ipynb)